

Reference

<https://www.youtube.com/watch?v=AwOlgOwaLl0>

Imports

```
import tensorflow_datasets as tfds
import tensorflow as tf
from tensorflow.keras.utils import to_categorical
```

Load Data

Explore more about dataset: https://www.tensorflow.org/datasets/catalog/tf_flowers

```
## Loading images and labels
(train_ds, train_labels), (test_ds, test_labels) = tfds.load("tf_flowers",
    split=["train[:70%]", "train[:30%]", ## Train test split
    batch_size=-1,
    as_supervised=True, # Include labels
)
Downloading and preparing dataset 218.21 MiB (download: 218.21 MiB,
generated: 221.83 MiB, total: 440.05 MiB) to
/rroot/tensorflow_datasets/tf_flowers/3.0.1...
{"model_id": "ffde15e341f544498356af997cb2caee", "version_major": 2, "version_minor": 0}

Dataset tf_flowers downloaded and prepared to
/rroot/tensorflow_datasets/tf_flowers/3.0.1. Subsequent calls will reuse this
data.
```

Image Preprocessing

```
## check existing image size
train_ds[0].shape
TensorShape([442, 1024, 3])

## Resizing images
train_ds = tf.image.resize(train_ds, (150, 150))
test_ds = tf.image.resize(test_ds, (150, 150))
train_ds[0].shape
TensorShape([150, 150, 3])

## Transforming labels to correct format
train_labels = to_categorical(train_labels, num_classes=5)
test_labels = to_categorical(test_labels, num_classes=5)
```

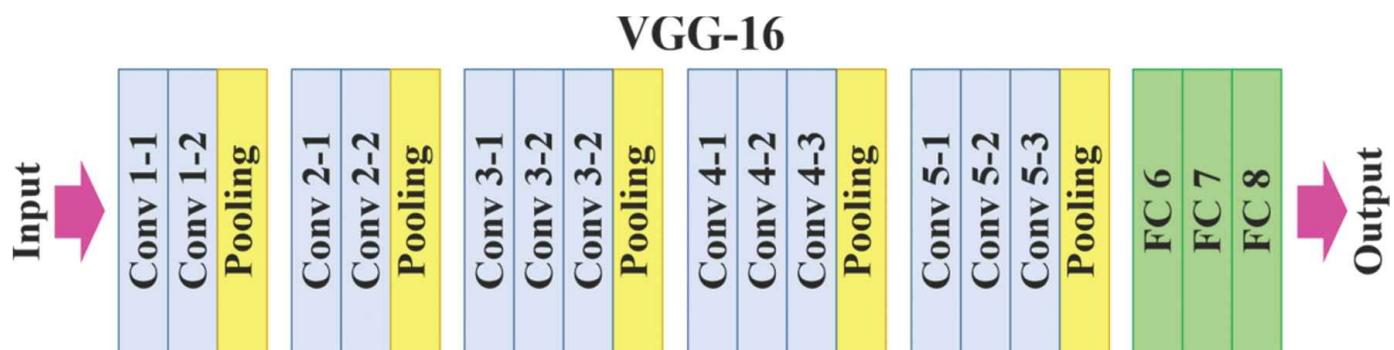
Use Pretrained VGG16 Image Classification model

Load a pre-trained CNN model trained on a large dataset

```
from tensorflow.keras.applications.vgg16 import VGG16
from tensorflow.keras.applications.vgg16 import preprocess_input

## Loading VGG16 model
base_model = VGG16(weights="imagenet", include_top=False,
input_shape=train_ds[0].shape)
```

Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/vgg16/vgg16_weights_tf_dim_ordering_tf_kernels_notop.h5 58889256/58889256 ————— 5s 0us/step



```
## We will not train base model i.e. Freeze Parameters in model's lower convolutional layers
base_model.trainable = False

## Preprocessing input
train_ds = preprocess_input(train_ds)
test_ds = preprocess_input(test_ds)

## model details
base_model.summary()

Model: "vgg16"
```

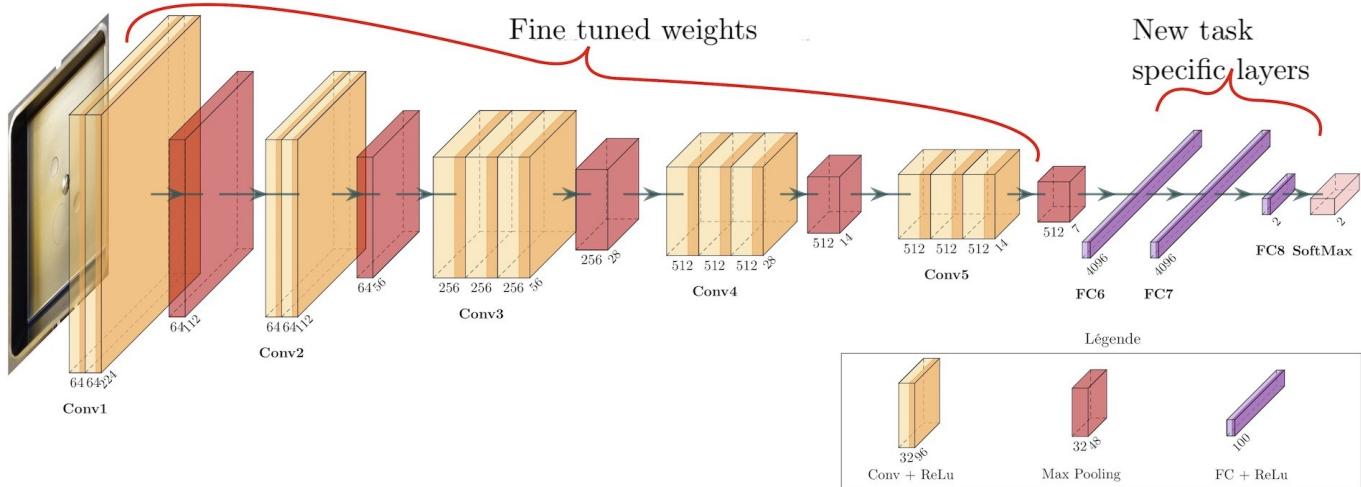
Layer (type) Param #	Output Shape
input_layer (InputLayer) 0	(None, 150, 150, 3)
block1_conv1 (Conv2D) 1,792	(None, 150, 150, 64)
block1_conv2 (Conv2D) 36,928	(None, 150, 150, 64)

	block1_pool (MaxPooling2D)	(None, 75, 75, 64)
0		
73,856	block2_conv1 (Conv2D)	(None, 75, 75, 128)
147,584	block2_conv2 (Conv2D)	(None, 75, 75, 128)
0	block2_pool (MaxPooling2D)	(None, 37, 37, 128)
295,168	block3_conv1 (Conv2D)	(None, 37, 37, 256)
590,080	block3_conv2 (Conv2D)	(None, 37, 37, 256)
590,080	block3_conv3 (Conv2D)	(None, 37, 37, 256)
0	block3_pool (MaxPooling2D)	(None, 18, 18, 256)
1,180,160	block4_conv1 (Conv2D)	(None, 18, 18, 512)
2,359,808	block4_conv2 (Conv2D)	(None, 18, 18, 512)
2,359,808	block4_conv3 (Conv2D)	(None, 18, 18, 512)
0	block4_pool (MaxPooling2D)	(None, 9, 9, 512)
2,359,808	block5_conv1 (Conv2D)	(None, 9, 9, 512)
2,359,808	block5_conv2 (Conv2D)	(None, 9, 9, 512)
	block5_conv3 (Conv2D)	(None, 9, 9, 512)

```

2,359,808 |
|-----|
| block5_pool (MaxPooling2D) | (None, 4, 4, 512) |
0 |-----|
|-----|
Total params: 14,714,688 (56.13 MB)
Trainable params: 0 (0.00 B)
Non-trainable params: 14,714,688 (56.13 MB)

```



Add custom classifier with two dense layers of trainable parameters to model

```

#add our layers on top of this model
from tensorflow.keras import layers, models

flatten_layer = layers.Flatten()
dense_layer_1 = layers.Dense(50, activation='relu')
dense_layer_2 = layers.Dense(20, activation='relu')
prediction_layer = layers.Dense(5, activation='softmax')

model = models.Sequential([
    base_model,
    flatten_layer,
    dense_layer_1,
    dense_layer_2,
    prediction_layer
])

```

Train classifier layers on training data available for task

```

model.compile(
    optimizer='adam',
    loss='categorical_crossentropy',
    metrics=['accuracy'],
)

history=model.fit(train_ds, train_labels, epochs=10, validation_split=0.2,
batch_size=32)

```

```

Epoch 1/5
65/65 33s 274ms/step - accuracy: 0.3618 - loss: 2.6283 -
val_accuracy: 0.5156 - val_loss: 1.1623
Epoch 2/5
65/65 21s 174ms/step - accuracy: 0.5954 - loss: 1.0192 -
val_accuracy: 0.6284 - val_loss: 1.0488
Epoch 3/5
65/65 11s 168ms/step - accuracy: 0.7250 - loss: 0.7869 -
val_accuracy: 0.6634 - val_loss: 1.0971
Epoch 4/5
65/65 21s 179ms/step - accuracy: 0.7849 - loss: 0.5597 -
val_accuracy: 0.6518 - val_loss: 0.9481
Epoch 5/5
65/65 8s 127ms/step - accuracy: 0.8399 - loss: 0.4183 -
val_accuracy: 0.7043 - val_loss: 0.9855

los, accurac=model.evaluate(test_ds,test_labels)
print("Loss: ",los,"Accuracy: ", accurac)

35/35 13s 385ms/step - accuracy: 0.8811 - loss: 0.3352
Loss:  0.3205247223377228 Accuracy:  0.8864668607711792

import matplotlib.pyplot as plt
plt.plot(history.history['accuracy'])
plt.title('ACCURACY')
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.legend(['train'],loc='upper left')
plt.show()

```

